applied-ml

Curated papers, articles, and blogs on **data science & machine learning in production**.

Figuring out how to implement your ML project? Learn how other organizations did it:

- **How** the problem is framed 🔎(e.g., personalization as recsys vs. search vs. sequences)
- **What** machine learning techniques worked ✔ (and sometimes, what didn't ❌)
- **Why** it works, the science behind it with research, literature, and references 👈
- **What** real-world results were achieved (so you can better assess ROI ⏱💰📈)

P.S., Want a summary of ML advancements? 👉 `ml-surveys`

Table of Contents

21. Audio
22. Validation and A/B Testing
23. Model Management
24. Efficiency
25. Ethics
26. Practices
27. Team Structure
28. Fails

Data Quality

1. Monitoring Data Quality at Scale with Statistical Modeling `Uber`
2. An Approach to Data Quality for Netflix Personalization Systems `Netflix`
3. Automating Large-Scale Data Quality Verification (Paper) `Amazon`
4. Meet Hodor — Gojek's Upstream Data Quality Tool `Gojek`
5. Reliable and Scalable Data Ingestion at Airbnb `Airbnb`
6. Data Management Challenges in Production Machine Learning (Paper) `Google`
7. Improving Accuracy By Certainty Estimation of Human Decisions, Labels, and Raters (Paper) `Facebook`

Data Engineering

1. Zipline: Airbnb's Machine Learning Data Management Platform `Airbnb`
2. Sputnik: Airbnb's Apache Spark Framework for Data Engineering `Airbnb`
3. Unbundling Data Science Workflows with Metaflow and AWS Step Functions `Netflix`
4. How DoorDash is Scaling its Data Platform to Delight Customers and Meet Growing Demand `DoorDash`
5. Revolutionizing Money Movements at Scale with Strong Data Consistency `Uber`

Data Discovery

1. Amundsen — Lyft's Data Discovery & Metadata Engine `Lyft`
2. Open Sourcing Amundsen: A Data Discovery And Metadata Platform (Code) `Lyft`
3. Amundsen: One Year Later `Lyft`
4. Using Amundsen to Support User Privacy via Metadata Collection at Square `Square`
5. Discovery and Consumption of Analytics Data at Twitter `Twitter`

6. Democratizing Data at Airbnb `Airbnb`
7. Databook: Turning Big Data into Knowledge with Metadata at Uber `Uber`
8. Turning Metadata Into Insights with Databook `Uber`
9. Metacat: Making Big Data Discoverable and Meaningful at Netflix (Code) `Netflix`
10. DataHub: A Generalized Metadata Search & Discovery Tool (Code) `LinkedIn`
11. DataHub: Popular Metadata Architectures Explained `LinkedIn`
12. How We Improved Data Discovery for Data Scientists at Spotify `Spotify`
13. How We're Solving Data Discovery Challenges at Shopify `Shopify`
14. Nemo: Data discovery at Facebook `Facebook`
15. Apache Atlas: Data Goverance and Metadata Framework for Hadoop (Code) `Apache`
16. Collect, Aggregate, and Visualize a Data Ecosystem's Metadata (Code) `WeWork`

Feature Stores

1. Introducing Feast: an open source feature store for machine learning (Code) `Gojek`
2. Feast: Bridging ML Models and Data `Gojek`
3. Building a Scalable ML Feature Store with Redis, Binary Serialization, and Compression `DoorDash`
4. Michelangelo Palette: A Feature Engineering Platform at Uber `Uber`
5. Distributed Time Travel for Feature Generation `Netflix`
6. The Architecture That Powers Twitter's Feature Store `Twitter`
7. Rapid Experimentation Through Standardization: Typed AI features for LinkedIn's Feed `LinkedIn`

Classification

1. High-Precision Phrase-Based Document Classification on a Modern Scale (Paper) `LinkedIn`
2. Chimera: Large-scale Classification using Machine Learning, Rules, and Crowdsourcing (Paper) `WalmartLabs`
3. Large-scale Item Categorization for e-Commerce (Paper) `DianPing, eBay`
4. Large-scale Item Categorization in e-Commerce Using Multiple Recurrent Neural Networks (Paper) `NAVER`
5. Categorizing Products at Scale `Shopify`
6. Learning to Diagnose with LSTM Recurrent Neural Networks (Paper) `Google`
7. Discovering and Classifying In-app Message Intent at Airbnb `Airbnb`
8. How We Built the Good First Issues Feature `GitHub`
9. Teaching Machines to Triage Firefox Bugs `Mozilla`
10. Testing Firefox More Efficiently with Machine Learning `Mozilla`

11. Using ML to Subtype Patients Receiving Digital Mental Health Interventions (Paper) `Microsoft`
12. Prediction of Advertiser Churn for Google AdWords (Paper) `Google`
13. Scalable Data Classification for Security and Privacy (Paper) `Facebook`

Regression

1. Using Machine Learning to Predict Value of Homes On Airbnb `Airbnb`
2. Using Machine Learning to Predict the Value of Ad Requests `Twitter`
3. Open-Sourcing Riskquant, a Library for Quantifying Risk (Code) `NetFlix`

Forecasting

1. Forecasting at Uber: An Introduction `Uber`
2. Engineering Extreme Event Forecasting at Uber with RNN `Uber`
3. Transforming Financial Forecasting with Data Science and Machine Learning at Uber `Uber`
4. Under the Hood of Gojek's Automated Forecasting Tool `GoJek`
5. BusTr: Predicting Bus Travel Times from Real-Time Traffic (Paper, Video) `Google`
6. Retraining Machine Learning Models in the Wake of COVID-19 `DoorDash`
7. Automatic Forecasting using Prophet, Databricks, Delta Lake and MLflow (Paper, Code) `Atlassian`

Recommendation

1. Amazon.com Recommendations: Item-to-Item Collaborative Filtering (Paper) `Amazon`
2. Temporal-Contextual Recommendation in Real-Time (Paper) `Amazon`
3. P-Companion: A Framework for Diversified Complementary Product Recommendation (Paper) `Amazon`
4. Recommending Complementary Products in E-Commerce Push Notifications (Paper) `Alibaba`
5. Behavior Sequence Transformer for E-commerce Recommendation in Alibaba (Paper) `Alibaba`
6. TPG-DNN: A Method for User Intent Prediction with Multi-task Learning (Paper) `Alibaba`
7. PURS: Personalized Unexpected Recommender System for Improving User Satisfaction (Paper) `Alibaba`
8. Session-based Recommendations with Recurrent Neural Networks (Paper) `Telefonica`
9. How 20th Century Fox uses ML to predict a movie audience (Paper) `20th Century Fox`
10. Deep Neural Networks for YouTube Recommendations `YouTube`
11. Personalized Recommendations for Experiences Using Deep Learning `TripAdvisor`
12. E-commerce in Your Inbox: Product Recommendations at Scale `Yahoo`

13. Product Recommendations at Scale (Paper) `Yahoo`
14. Powered by AI: Instagram's Explore recommender system `Facebook`
15. Netflix Recommendations: Beyond the 5 stars (Part 1 (Part 2) `Netflix`
16. Learning a Personalized Homepage `Netflix`
17. Artwork Personalization at Netflix `Netflix`
18. To Be Continued: Helping you find shows to continue watching on Netflix `Netflix`
19. Calibrated Recommendations (Paper) `Netflix`
20. Food Discovery with Uber Eats: Recommending for the Marketplace `Uber`
21. Food Discovery with Uber Eats: Using Graph Learning to Power Recommendations `Uber`
22. How Music Recommendation Works — And Doesn't Work `Spotify`
23. Music recommendation at Spotify `Spotify`
24. Recommending Music on Spotify with Deep Learning `Spotify`
25. For Your Ears Only: Personalizing Spotify Home with Machine Learning `Spotify`
26. Reach for the Top: How Spotify Built Shortcuts in Just Six Months `Spotify`
27. Explore, Exploit, and Explain: Personalizing Explainable Recommendations with Bandits (Paper) `Spotify`
28. Contextual and Sequential User Embeddings for Large-Scale Music Recommendation (Paper) `Spotify`
29. The Evolution of Kit: Automating Marketing Using Machine Learning `Shopify`
30. Using Machine Learning to Predict what File you Need Next (Part 1) `Dropbox`
31. Using Machine Learning to Predict what File you Need Next (Part 2) `Dropbox`
32. Personalized Recommendations in LinkedIn Learning `LinkedIn`
33. A Closer Look at the AI Behind Course Recommendations on LinkedIn Learning (Part 1) `LinkedIn`
34. A Closer Look at the AI Behind Course Recommendations on LinkedIn Learning (Part 2) `LinkedIn`
35. Learning to be Relevant: Evolution of a Course Recommendation System (**PAPER NEEDED**) `LinkedIn`
36. Building a Heterogeneous Social Network Recommendation System `LinkedIn`
37. How TikTok recommends videos #ForYou `ByteDance`
38. A Meta-Learning Perspective on Cold-Start Recommendations for Items (Paper) `Twitter`
39. Zero-Shot Heterogeneous Transfer Learning from RecSys to Cold-Start Search Retrieval (Paper) `Google`
40. Improved Deep & Cross Network for Feature Cross Learning in Web-scale LTR Systems (Paper) `Google`
41. Personalized Channel Recommendations in Slack `Slack`
42. Deep Retrieval: End-to-End Learnable Structure Model for Large-Scale Recommendations (Paper) `ByteDance`

43. Future Data Helps Training: Modeling Future Contexts for Session-based Recommendation (Paper) `Tencent`
44. Using AI to Help Health Experts Address the COVID-19 Pandemic `Facebook`
45. A Case Study of Session-based Recommendations in the Home-improvement Domain (Paper) `Home Depot`
46. Balancing Relevance and Discovery to Inspire Customers in the IKEA App (Paper) `Ikea`
47. How we use AutoML, Multi-task learning and Multi-tower models for Pinterest Ads `Pinterest`

Search & Ranking

1. Amazon Search: The Joy of Ranking Products (Paper, Video, Code) `Amazon`
2. Why Do People Buy Seemingly Irrelevant Items in Voice Product Search? (Paper) `Amazon`
3. How Lazada Ranks Products to Improve Customer Experience and Conversion `Lazada`
4. Using Deep Learning at Scale in Twitter's Timelines `Twitter`
5. Machine Learning-Powered Search Ranking of Airbnb Experiences `Airbnb`
6. Applying Deep Learning To Airbnb Search (Paper) `Airbnb`
7. Managing Diversity in Airbnb Search (Paper) `Airbnb`
8. Improving Deep Learning for Airbnb Search (Paper) `Airbnb`
9. Ranking Relevance in Yahoo Search (Paper) `Yahoo`
10. An Ensemble-based Approach to Click-Through Rate Prediction for Promoted Listings at Etsy (Paper) `Etsy`
11. Learning to Rank Personalized Search Results in Professional Networks (Paper) `LinkedIn`
12. Entity Personalized Talent Search Models with Tree Interaction Features (Paper) `LinkedIn`
13. In-session Personalization for Talent Search (Paper) `LinkedIn`
14. The AI Behind LinkedIn Recruiter search and recommendation systems `LinkedIn`
15. Quality Matches Via Personalized AI for Hirer and Seeker Preferences `LinkedIn`
16. Understanding Dwell Time to Improve LinkedIn Feed Ranking `LinkedIn`
17. Ads Allocation in Feed via Constrained Optimization (Paper, Video) `LinkedIn`
18. AI at Scale in Bing `Microsoft`
19. Query Understanding Engine in Traveloka Universal Search `Traveloka`
20. The Secret Sauce Behind Search Personalisation `GoJek`
21. Food Discovery with Uber Eats: Building a Query Understanding Engine `Uber`
22. Neural Code Search: ML-based Code Search Using Natural Language Queries `Facebook`
23. Bayesian Product Ranking at Wayfair `Wayfair`
24. COLD: Towards the Next Generation of Pre-Ranking System (Paper) `Alibaba`

25. Understanding Searches Better Than Ever Before (Paper) `Google`
26. Shop The Look: Building a Large Scale Visual Shopping System at Pinterest (Paper, Video) `Pinterest`
27. GDMix: A Deep Ranking Personalization Framework (Code) `LinkedIn`
28. Bringing Personalized Search to Etsy `Etsy`
29. Building a Better Search Engine for Semantic Scholar `Allen Institute for AI`

Embeddings

1. Billion-scale Commodity Embedding for E-commerce Recommendation in Alibaba (Paper) `Alibaba`
2. Embeddings@Twitter `Twitter`
3. Listing Embeddings in Search Ranking (Paper) `Airbnb`
4. Understanding Latent Style `Stitch Fix`
5. Towards Deep and Representation Learning for Talent Search at LinkedIn (Paper) `LinkedIn`
6. Vector Representation Of Items, Customer And Cart To Build A Recommendation System (Paper) `Sears`
7. Machine Learning for a Better Developer Experience `Netflix`
8. Announcing ScaNN: Efficient Vector Similarity Search (Paper, Code) `Google`

Natural Language Processing

1. Abusive Language Detection in Online User Content (Paper) `Yahoo`
2. How Natural Language Processing Helps LinkedIn Members Get Support Easily `LinkedIn`
3. Building Smart Replies for Member Messages `LinkedIn`
4. DeText: A deep NLP Framework for Intelligent Text Understanding (Code) `LinkedIn`
5. Smart Reply: Automated Response Suggestion for Email (Paper) `Google`
6. Gmail Smart Compose: Real-Time Assisted Writing (Paper) `Google`
7. SmartReply for YouTube Creators `Google`
8. Using Neural Networks to Find Answers in Tables (Paper) `Google`
9. A Scalable Approach to Reducing Gender Bias in Google Translate `Google`
10. Assistive AI Makes Replying Easier `Microsoft`
11. AI Advances to Better Detect Hate Speech `Facebook`
12. A State-of-the-Art Open Source Chatbot (Paper) `Facebook`
13. A Highly Efficient, Real-Time Text-to-Speech System Deployed on CPUs `Facebook`
14. Deep Learning to Translate Between Programming Languages (Paper, Code) `Facebook`
15. Deploying Lifelong Open-Domain Dialogue Learning (Paper) `Facebook`

16. Goal-Oriented End-to-End Conversational Models with Profile Features in a Real-World Setting (Paper) `Amazon`
17. How Gojek Uses NLP to Name Pickup Locations at Scale `GoJek`
18. Give Me Jeans not Shoes: How BERT Helps Us Deliver What Clients Want `Stitch Fix`
19. The State-of-the-art Open-Domain Chatbot in Chinese and English (Paper) `Baidu`
20. PEGASUS: A State-of-the-Art Model for Abstractive Text Summarization (Paper, Code) `Google`
21. Photon: A Robust Cross-Domain Text-to-SQL System (Paper) (Demo) `Salesforce`
22. GeDi: A Powerful New Method for Controlling Language Models (Paper, Code) `Salesforce`
23. Applying Topic Modeling to Improve Call Center Operations `RICOH`
24. WIDeText: A Multimodal Deep Learning Framework `Airbnb`

Sequence Modelling

1. Practice on Long Sequential User Behavior Modeling for Click-Through Rate Prediction (Paper) `Alibaba`
2. Search-based User Interest Modeling with Sequential Behavior Data for CTR Prediction (Paper) `Alibaba`
3. Deep Learning for Electronic Health Records (Paper) `Google`
4. Deep Learning for Understanding Consumer Histories (Paper) `Zalando`
5. Continual Prediction of Notification Attendance with Classical and Deep Networks (Paper) `Telefonica`
6. Using Recurrent Neural Network Models for Early Detection of Heart Failure Onset (Paper) `Sutter Health`
7. Doctor AI: Predicting Clinical Events via Recurrent Neural Networks (Paper) `Sutter Health`
8. How Duolingo uses AI in every part of its app `Duolingo`
9. Leveraging Online Social Interactions For Enhancing Integrity at Facebook (Paper, Video) `Facebook`

Computer Vision

1. Categorizing Listing Photos at Airbnb `Airbnb`
2. Amenity Detection and Beyond — New Frontiers of Computer Vision at Airbnb `Airbnb`
3. Powered by AI: Advancing product understanding and building new shopping experiences `Facebook`
4. Creating a Modern OCR Pipeline Using Computer Vision and Deep Learning `Dropbox`
5. How we Improved Computer Vision Metrics by More Than 5% Only by Cleaning Labelling Errors `Deepomatic`
6. A Neural Weather Model for Eight-Hour Precipitation Forecasting (Paper) `Google`

7. Machine Learning-based Damage Assessment for Disaster Relief (Paper) `Google`
8. RepNet: Counting Repetitions in Videos (Paper) `Google`
9. Converting Text to Images for Product Discovery (Paper) `Amazon`
10. How Disney Uses PyTorch for Animated Character Recognition `Disney`
11. Image Captioning as an Assistive Technology (Video) `IBM`
12. AI for AG: Production machine learning for agriculture `Blue River`
13. AI for Full-Self Driving at Tesla `Tesla`
14. On-device Supermarket Product Recognition `Google`
15. Using Machine Learning to Detect Deficient Coverage in Colonoscopy Screenings (Paper) `Google`
16. Shop The Look: Building a Large Scale Visual Shopping System at Pinterest (Paper, Video) `Pinterest`
17. Developing Real-Time, Automatic Sign Language Detection for Video Conferencing (Paper) `Google`

Reinforcement Learning

1. Deep Reinforcement Learning for Sponsored Search Real-time Bidding (Paper) `Alibaba`
2. Dynamic Pricing on E-commerce Platform with Deep Reinforcement Learning (Paper) `Alibaba`
3. Budget Constrained Bidding by Model-free Reinforcement Learning in Display Advertising (Paper) `Alibaba`
4. Productionizing Deep Reinforcement Learning with Spark and MLflow `Zynga`
5. Deep Reinforcement Learning in Production Part1 Part 2 `Zynga`
6. Building AI Trading Systems `Denny Britz`

Anomaly Detection

1. Detecting Performance Anomalies in External Firmware Deployments `Netflix`
2. Detecting and Preventing Abuse on LinkedIn using Isolation Forests (Code) `LinkedIn`
3. Preventing Abuse Using Unsupervised Learning `LinkedIn`
4. The Technology Behind Fighting Harassment on LinkedIn `LinkedIn`
5. Uncovering Insurance Fraud Conspiracy with Network Learning (Paper) `Ant Financial`
6. How Does Spam Protection Work on Stack Exchange? `Stack Exchange`
7. Auto Content Moderation in C2C e-Commerce `Mercari`
8. Blocking Slack Invite Spam With Machine Learning `Slack`
9. Cloudflare Bot Management: Machine Learning and More `Cloudflare`
10. Anomalies in Oil Temperature Variations in a Tunnel Boring Machine `SENER`
11. Using Anomaly Detection to Monitor Low-Risk Bank Customers `Rabobank`

12. Fighting fraud with Triplet Loss `OLX Group`
13. Facebook is Now Using AI to Sort Content for Quicker Moderation (Alternative) `Facebook`
14. How AI is getting better at detecting hate speech Part 1, Part 2, Part 3, Part 4 `Facebook`
15. Deep Anomaly Detection with Spark and Tensorflow (Hopsworks Video) `Swedbank , Hopsworks`

Graph

1. Building The LinkedIn Knowledge Graph `LinkedIn`
2. Retail Graph — Walmart's Product Knowledge Graph `Walmart`
3. Food Discovery with Uber Eats: Using Graph Learning to Power Recommendations `Uber`
4. AliGraph: A Comprehensive Graph Neural Network Platform (Paper) `Alibaba`
5. Scaling Knowledge Access and Retrieval at Airbnb `Airbnb`
6. Traffic Prediction with Advanced Graph Neural Networks `DeepMind`
7. SimClusters: Community-Based Representations for Recommendations (Paper, Video) `Twitter`

Optimization

1. How Trip Inferences and Machine Learning Optimize Delivery Times on Uber Eats `Uber`
2. Next-Generation Optimization for Dasher Dispatch at DoorDash `DoorDash`
3. Matchmaking in Lyft Line (Part 1) (Part 2) (Part 3) `Lyft`
4. The Data and Science behind GrabShare Carpooling (**PAPER NEEDED**) `Grab`
5. Optimization of Passengers Waiting Time in Elevators Using Machine Learning `Thyssen Krupp AG`
6. Think out of the package: Recommending package types for e-commerce shipments (Paper) `Amazon`

Information Extraction

1. Unsupervised Extraction of Attributes and Their Values from Product Description (Paper) `Rakuten`
2. Information Extraction from Receipts with Graph Convolutional Networks `Nanonets`
3. Using Machine Learning to Index Text from Billions of Images `Dropbox`
4. Extracting Structured Data from Templatic Documents (Paper) `Google`
5. AutoKnow: self-driving knowledge collection for products of thousands of types (Paper, Video) `Amazon`
6. One-shot Text Labeling using Attention and Belief Propagation for Information Extraction (Paper) `Alibaba`

Weak Supervision

1. Snorkel DryBell: A Case Study in Deploying Weak Supervision at Industrial Scale (Paper) `Google`
2. Osprey: Weak Supervision of Imbalanced Extraction Problems without Code (Paper) `Intel`
3. Overton: A Data System for Monitoring and Improving Machine-Learned Products (Paper) `Apple`
4. Bootstrapping Conversational Agents with Weak Supervision (Paper) `IBM`

Generation

1. Better Language Models and Their Implications (Paper) `OpenAI`
2. Language Models are Few-Shot Learners (Paper) (GPT-3 Blog post) `OpenAI`
3. Image GPT (Paper, Code) `OpenAI`
4. Deep Learned Super Resolution for Feature Film Production (Paper) `Pixar`
5. Unit Test Case Generation with Transformers `Microsoft`

Audio

1. Improving On-Device Speech Recognition with VoiceFilter-Lite (Paper) `Google`
2. The Machine Learning Behind Hum to Search `Google`

Validation and A/B Testing

1. The Reusable Holdout: Preserving Validity in Adaptive Data Analysis (Paper) `Google`
2. Detecting Interference: An A/B Test of A/B Tests `LinkedIn`
3. Experimenting to Solve Cramming `Twitter`
4. Announcing a New Framework for Designing Optimal Experiments with Pyro (Paper) (Paper) `Uber`
5. Enabling 10x More Experiments with Traveloka Experiment Platform `Traveloka`
6. Large Scale Experimentation at Stitch Fix (Paper) `Stitch Fix`
7. Multi-Armed Bandits and the Stitch Fix Experimentation Platform `Stitch Fix`
8. Experimentation with Resource Constraints `Stitch Fix`
9. Modeling Conversion Rates and Saving Millions Using Kaplan-Meier and Gamma Distributions (Code) `Better`
10. Computational Causal Inference at Netflix (Paper) `Netflix`
11. Key Challenges with Quasi Experiments at Netflix `Netflix`
12. Constrained Bayesian Optimization with Noisy Experiments (Paper) `Facebook`
13. Supporting Rapid Product Iteration with an Experimentation Analysis Platform `Curie`

14. Our Evolution Towards T-REX: The Prehistory of Experimentation Infrastructure at LinkedIn `LinkedIn`
15. How to Use Quasi-experiments and Counterfactuals to Build Great Products `Shopify`
16. Improving Online Experiment Capacity by 4X with Parallelization and Increased Sensitivity `DoorDash`
17. Spotify's New Experimentation Platform (Part 1) (Part 2) `Spotify`
18. Overlapping Experiment Infrastructure: More, Better, Faster Experimentation (Paper) `Google`

Model Management

1. Runway - Model Lifecycle Management at Netflix `Netflix`

Efficiency

1. GrokNet: Unified Computer Vision Model Trunk and Embeddings For Commerce (Paper) `Facebook`
2. Permute, Quantize, and Fine-tune: Efficient Compression of Neural Networks (Paper) `Uber`

Ethics

1. Building Inclusive Products Through A/B Testing (Paper) `LinkedIn`
2. LiFT: A Scalable Framework for Measuring Fairness in ML Applications (Paper) `LinkedIn`

Practices

1. Practical Recommendations for Gradient-Based Training of Deep Architectures (Paper) `Yoshua Bengio`
2. Machine Learning: The High Interest Credit Card of Technical Debt (Paper) (Paper) `Google`
3. Rules of Machine Learning: Best Practices for ML Engineering `Google`
4. On Challenges in Machine Learning Model Management `Amazon`
5. Machine Learning in Production: The Booking.com Approach `Booking`
6. 150 Successful Machine Learning Models: 6 Lessons Learned at Booking.com (Paper) `Booking`
7. Successes and Challenges in Adopting Machine Learning at Scale at a Global Bank `Rabobank`

Team structure

1. Engineers Shouldn't Write ETL: A Guide to Building a High Functioning Data Science Department `Stitch Fix`

2. Beware the Data Science Pin Factory: The Power of the Full-Stack Data Science Generalist `Stitch Fix`
3. Cultivating Algorithms: How We Grow Data Science at Stitch Fix `StitchFix`
4. Analytics at Netflix: Who We Are and What We Do `Netflix`

Fails

1. 160k+ High School Students Will Graduate Only If a Model Allows Them to `International Baccalaureate`
2. When It Comes to Gorillas, Google Photos Remains Blind `Google`
3. An Algorithm That 'Predicts' Criminality Based on a Face Sparks a Furor `Harrisburg University`
4. It's Hard to Generate Neural Text From GPT-3 About Muslims `OpenAI`
5. A British AI Tool to Predict Violent Crime Is Too Flawed to Use `United Kingdom`
6. More in awful-ai

P.S., Want a summary of ML advancements? Get up to speed with survey papers 👉 `ml-surveys`